INTERNATIONAL INSTITUTE OF MEDICAL SCIENCE & TECHNOLOGY COUNCIL

A Mini Project Documentation on

"Healthcare Cost Predictor"

Submitted in partial fulfillment of the requirement for internship on

"AIML and Gen AI"

Submitted by

COHORT _ 03

Bushra Yasmeen

Awaiz kha

Vaishali L N

Ashith HK

Shreeya D

Mahalakshmi DS

Pallavi Badiger

Mohammad sulaiman. H

Mohammed Inayathulla

Ikshu Dhanva H M

Safeya Zawath Zakir

Aishwarya B R

Sumith Shanthinath Chougale

Harshit Vaidya

Kallamadi BanuPrakash Reddy

Table Of Contents

SL No.	Title	Page No.
1.	Introduction	1
2.	Overview	2
3.	SDG Alignment	2-3
4.	Why SDG 3 Matters?	3
5.	Healthcare Cost Constraints	3-4
6.	Methodology	4-7
7.	System Workflow	7
8.	Data Source	8-11
9.	Variables Used	11-12
10.	Dummy Variable Explanation	12-13
11.	GitHub Implementation	13
12.	Model Explanation	13-14
13.	Is Unexplained GAP Present?	15
14.	Code Snippet	15-17
15.	Libraries Used and Explanation	17-19
16.	Conclusion	22-23
17.	Limitations	19
18.	Challenges Faced and Solutions	19-20

19.	Ethical Considerations	20-21
20.	Future Enhancements	21
21.	Industry Use Cases	21-22
22.	References	23

HEALTHCARE COST PREDICTION MODEL

INTRODUCTION

Healthcare expenses are increasing globally due to rising medical costs, lifestyle-related diseases, and growing healthcare demands. Medical insurance plays a critical role in managing this financial uncertainty by protecting individuals and families from unexpected medical burdens. However, insurance premiums often vary significantly based on personal attributes such as age, body mass index (BMI), smoking status, and other demographic factors. Many individuals are unaware of how strongly these lifestyle and health conditions influence their insurance charges. Therefore, predicting insurance costs based on individual health characteristics can help people better understand the financial impact of their lifestyle decisions and encourage preventive healthcare practices.

This project focuses on building a Healthcare Cost Predictor using machine learning techniques to estimate annual medical insurance charges. The system begins with data preprocessing steps such as handling missing values, removing duplicates, detecting outliers using the IQR method, and encoding categorical variables into numerical form. After cleaning and preparing the dataset, correlation analysis and heatmap visualizations are performed to identify the most influential factors affecting insurance costs. Special emphasis is placed on evaluating the impact of smoking habits, BMI, and age, as these variables are closely linked to long-term health risks and medical expenditures.

Regression-based machine learning models, including Linear Regression, Decision Tree, and Random Forest, are trained and evaluated using performance metrics such as MAE, RMSE, and R^2 score. By comparing these models, the most accurate predictive model is selected to provide reliable insurance cost estimations. The project not only demonstrates the practical application of data analysis and regression techniques but also aligns with Sustainable Development Goal 3 (Good Health and Well-Being) by promoting awareness of how lifestyle factors influence healthcare costs. Through data-driven insights, the system encourages informed decision making understanding of health-related financial planning.

OVERVIEW

Healthcare Cost Predictor refers to the use of artificial intelligence and machine learning techniques to estimate medical insurance charges based on individual health and demographic characteristics. By leveraging regression models and data analysis methods, this system analyses factors such as age, BMI, smoking status, gender, and region to generate accurate insurance cost predictions. The system helps identify how lifestyle choices and health conditions influence annual medical expenses. It supports data-driven decision-making for individuals and promotes awareness about preventive healthcare practices.

Key aspects of the Healthcare Cost Predictor include:

- Data Collection and Processing
- Data Cleaning and Outlier Detection (IQR Method)
- Categorical Variable Encoding (Label/Dummy Encoding)
- Correlation Analysis and Heatmap Visualization
- Feature Selection and Impact Analysis
- Machine Learning Model Training (Linear Regression, Decision Tree, Random Forest)
- Model Evaluation (MAE, RMSE, R^2 Score)
- Prediction and Visualization of Results

This structured approach ensures accurate cost estimation while highlighting the most influential health-related factors, such as smoking and BMI, that significantly impact insurance premiums.

SDG ALIGNMENT

This project aligns with **United Nations Sustainable Development Goal 3: Good Health and Wellbeing**, specifically target 3.8, which aims to *"Achieve universal health coverage, including financial risk protection"*. By providing a predictive model for insurance costs, the Health Cost Predictor helps individuals and insurers anticipate and plan for medical expenses. Improved cost forecasting enables more affordable and fair healthcare financing, aiding universal health coverage and access to quality medical services. For example, if individuals can estimate their future premiums, they can budget accordingly, reducing financial strain and

the risk of unexpected healthcare costs. Thus, the project contributes to making healthcare more predictable and accessible, supporting SDG 3's emphasis on affordable health coverage. Healthcare cost prediction is commonly calculated using certain measurable constraints (input variables) that significantly influence annual medical insurance charges. These constraints represent demographic, physical, and lifestyle-related characteristics of an individual. The prediction model uses these factors as independent variables to estimate the dependent variable, which is the insurance charge.

WHY SDG 3 MATTERS?

Sustainable Development Goal 3 aims to ensure healthy lives and promote well-being for all at all ages, making it one of the most essential global development goals. Good health is the foundation for education, employment, economic growth, and overall quality of life. Without access to affordable healthcare and preventive awareness, individuals and families may face severe financial and social challenges. SDG 3 focuses not only on treating diseases but also on preventing them by encouraging healthy lifestyles, reducing risk factors such as smoking and obesity, and strengthening healthcare systems. In the context of our Healthcare Cost Predictor project, this goal is highly relevant because the project analyses how factors like age, BMI, and smoking status influence medical insurance costs. By using data and machine learning to highlight the financial impact of health-related behaviours, the project promotes awareness, supports informed decision-making, and contributes to the broader mission of improving health and well-being in society.

HEALTHCARE COST CONSTRAINTS

The major constraints used in this project include:

- **Age** – Medical risk generally increases with age, leading to higher insurance premiums.
- **BMI (Body Mass Index)** – Higher BMI may indicate obesity-related health risks, increasing healthcare expenses.
- **Smoker Status** – Smoking significantly increases the probability of severe diseases, resulting in higher insurance charges.
- **Sex** – Gender may influence healthcare cost patterns.
- **Region** – Insurance costs may vary slightly based on geographical location.

The general regression-based prediction formula can be represented as:

Insurance Charges = $\beta_0 + \beta_1$ (Age) + β_2 (BMI) + β_3 (Smoker) + β_4 (Sex) + β_5 (Region) + ε

Where:

β_0 = Intercept β_1, β_2, β_3, β_4, β_5 = Coefficients representing impact of each constraint ε = Error term

Example

Consider the following individual:

- Age = 40 years
- BMI = 32
- Smoker = Yes (1)
- Sex = Male (1)
- Region = Southeast

Since the person is a smoker and has a high BMI, the model will assign higher coefficient contributions to these variables, resulting in a higher predicted insurance charge compared to a non-smoker with lower BMI.

Predicted Insurance Cost ≈ ₹45,000 – ₹60,000 annually (depending on trained model output)
Whereas a 25-year-old non-smoker with BMI 22 may have:

Predicted Insurance Cost ≈ ₹12,000 – ₹18,000 annually.

This demonstrates how the constraints directly influence the prediction outcome and why smoking and BMI are often found to be the most impactful features in healthcare cost analysis.

METHODOLOGY

The methodology of this project follows a systematic machine learning pipeline to ensure accurate prediction of medical insurance charges. The process includes data collection, preprocessing, feature analysis, model building, evaluation, and interpretation of results.

1. Data Collection

The dataset was obtained from the Kaggle repository *Medical Cost Personal Dataset*. It contains 1,338 records with demographic and health-related information such as age, BMI, smoking status, and insurance charges. The dataset is publicly available and anonymized.

2. Data Pre-processing

Data pre-processing was performed to prepare the dataset for machine learning.

- *Missing Value Check:* The dataset was examined using Pandas functions, and no missing values were found.
- *Duplicate Removal:* Duplicate rows were identified and removed to maintain data consistency.
- *Outlier Detection:* Outliers were analysed using summary statistics and distribution analysis. Since medical costs naturally show variation, extreme values were retained to preserve realistic variability.
- *Categorical Encoding:* Categorical variables such as **sex**, **smoker**, and **region** were converted into numerical form using Label Encoding and Dummy Variable techniques. This transformation ensures compatibility with regression models.

3. Exploratory Data Analysis (EDA)

Exploratory Data Analysis was conducted to understand relationships among features.

- Generated summary statistics (mean, median, quartiles)
- Examined distribution of charges (right-skewed)
- Created correlation matrix and heatmap
- Identified key predictors (smoker, age, BMI)

EDA helped in selecting relevant features for model training.

4. Feature Selection

Based on correlation analysis and domain knowledge, the following features were selected for model building:

- Age
- BMI
- Smoker

- Sex
- Region

Among these, smoker, age, and BMI showed the strongest influence on insurance charges.

5. Model Development

Since the target variable (charges) is continuous, a **Linear Regression model** was chosen as the primary algorithm.

Steps involved:

1. Splitting dataset into training (80%) and testing (20%) sets
2. Training the Linear Regression model using training data
3. Learning regression coefficients
4. Generating predictions on test dataset

Linear regression was selected due to its simplicity, interpretability, and ability to model linear relationships.

6. Model Evaluation

The model performance was evaluated using:

- **R^2 Score** – Measures proportion of variance explained
- **Mean Squared Error (MSE)** – Measures squared prediction error
- **Root Mean Squared Error (RMSE)** – Interpretable prediction error

Similar R^2 scores for training and testing sets indicate good generalization ability.

7. Result Interpretation

The trained model revealed that:

- Smoking status has the highest impact on insurance charges
- Age has a positive correlation with medical cost
- BMI moderately influences cost
- Other variables have relatively smaller effects

The findings align with real-world healthcare trends.

8. Conclusion of Methodology

The methodology follows a structured data science pipeline:

Data Collection → Data Cleaning → Encoding → Feature Analysis → Model Training → Evaluation → Interpretation

This systematic approach ensures reliable predictions and meaningful insights into healthcare cost determinants.

SYSTEM WORKFLOW

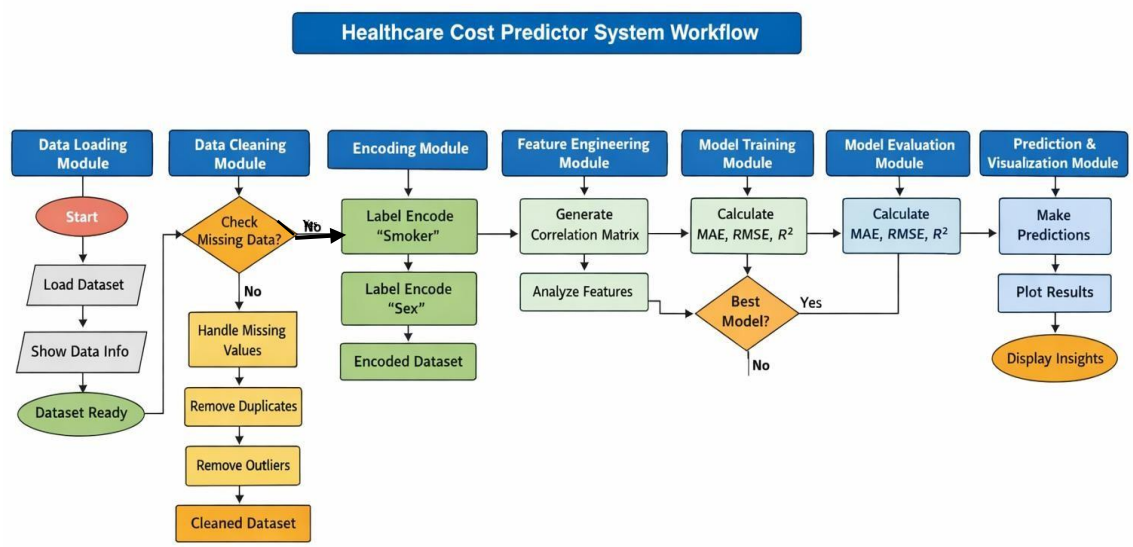

Figure 1: System Workflow of the Proposed System

The Figure 1 represents the complete workflow of the Healthcare Cost Predictor System. It begins with the Data Loading Module, where the dataset is imported and basic information is examined. In the Data Cleaning Module, missing values are checked, duplicates are removed, and outliers are handled to prepare a clean dataset. The Encoding Module converts categorical variables like smoker and sex into numerical format for model compatibility.

Next, the Feature Engineering Module generates a correlation matrix and analyses key features affecting insurance charges. The Model Training Module trains regression models and calculates performance metrics. The Model Evaluation Module compares models using MAE, RMSE, and R^2 to select the best-performing model. Finally, the Prediction and Visualization Module generates predictions, plots results, and displays insights regarding healthcare cost determinants.

DATA SOURCE

- Data Acquisition and Source Authentication

The dataset used in this project is the public *Medical Insurance Personal Dataset*, commonly referred to as **insurance.csv**, which is widely used for predictive modelling practice and regression-based tutorials. The dataset originates from the Kaggle repository titled *"Medical Cost Personal Dataset."* The CSV file was downloaded from Kaggle and contains 1,338 individual records of insurance-related data. Since the dataset is publicly available, it is freely accessible and suitable for academic and research purposes. The authenticity of the dataset was verified by referencing the official Kaggle source and related documentation. Furthermore, the dataset does not contain any personally identifiable information (PII), as it includes only anonymized demographic and health-related attributes, making it safe for analysis and educational use.

- Dataset Description

The dataset consists of 1,338 entries and 7 columns, where each row represents an individual's insurance-related information. The features included in the dataset are age, sex, BMI, children, smoker, region, and charges. The variable *age* represents the age of the primary beneficiary in years, while *sex* indicates gender (male or female). The *BMI* variable represents Body Mass Index measured in kg/m². The *children* variable denotes the number of dependents covered under the insurance plan. The *smoker* variable identifies whether the individual is a smoker or non-smoker, and *region* specifies the residential area in the United States (northeast, southeast, southwest, northwest). The *charges* variable represents the annual medical insurance cost in USD and serves as the target variable for prediction. Summary statistics show that age ranges from 18 to 64 years, with a mean of approximately 39 years. BMI values range from about 16 to 53, with a mean of 30.66, indicating that many individuals fall into the overweight category. Insurance charges range from approximately $1,122 to $63,770. The mean charges are around $13,270, while the median is approximately $9,382. Since the mean is higher than the median, the distribution of charges is right-skewed, suggesting the presence of high-cost outliers. Smoking status appears to be a critical contributing factor, as smokers tend to have significantly higher insurance charges compared to non-smokers.

DATA PRE-PROCESSING

The pre-processing phase began with verifying the presence of missing values. Using Pandas, it was confirmed that none of the 1,338 records contained missing data. Each column had complete entries, eliminating the need for imputation. This clean dataset simplified model preparation. Following this, summary statistics revealed that the *charges* variable is heavily skewed, indicating the presence of extreme high-cost individuals. Although techniques such as log transformation or trimming extreme values could have been applied, the full dataset was retained to preserve real-world medical cost variability. Next, categorical variables were converted into numerical form to ensure compatibility with regression models. The smoker variable was encoded as yes = 1 and no = 0, while sex was encoded as male = 1 and female = 0. The region variable was encoded using label or dummy encoding techniques. This transformation allowed the regression model to process categorical information effectively. A correlation matrix and heatmap were generated to analyse relationships among variables. The results indicated that smoking status has a strong positive correlation with insurance charges, while age and BMI show moderate positive correlations. In contrast, sex, number of children, and region demonstrate weaker correlations. This analysis guided feature selection for model building.

FEATURE ANALYSIS

Feature analysis was conducted to examine the individual impact of each variable on insurance charges. It was observed that insurance charges tend to increase with age, as older individuals generally have higher medical risk. BMI was also found to influence medical costs, since higher BMI values are associated with increased health complications. Smoking status emerged as the most significant predictor, with smokers incurring substantially higher insurance charges compared to non-smokers. The influence of sex was relatively minor compared to lifestyle-related factors, while the number of children showed minimal direct impact on personal medical costs. Regional differences existed but were not dominant contributors. Overall, smoking status, age, and BMI emerged as the strongest predictors of insurance charges.

MODEL SELECTION AND JUSTIFICATION

Linear Regression was selected as the primary model because the target variable, insurance charges, is continuous in nature. A linear relationship was hypothesized between the independent variables (age, BMI, smoker, etc.) and the dependent variable. Linear regression provides high interpretability, as each coefficient quantifies the impact of a corresponding feature on predicted insurance cost. Although linear regression assumes linearity and can be sensitive to outliers, it was considered an appropriate baseline model due to dataset simplicity and the need for interpretability in healthcare applications.

MODEL IMPLEMENTATION

The model was implemented using Python and the Scikit-learn library. The dataset was divided into training and testing subsets using an 80/20 split. The Linear Regression model was instantiated and trained on the training dataset. After training, predictions were generated on the test dataset. The regression coefficients were extracted to interpret the contribution of each feature to the predicted insurance charges.

TRAIN-TEST SPLIT JUSTIFICATION

The dataset was divided into training and testing subsets using an 80/20 ratio to evaluate the model on unseen data. This approach prevents overfitting and ensures that the model generalizes well beyond the training dataset. Separating the data allows for reliable performance estimation and provides a realistic evaluation of prediction accuracy.

MODEL EVALUATION

The model was evaluated using two primary performance metrics: R^2 (Coefficient of Determination) and Mean Squared Error (MSE). The R^2 score measures the proportion of variance in insurance charges explained by the model. The training R^2 was approximately 0.748, and the testing R^2 was also approximately 0.748. This indicates that about 74.8% of the variability in insurance charges is explained by the model and that the model generalizes well to unseen data. Mean Squared Error measures the average squared difference between predicted and actual values, while RMSE provides the average prediction error in dollar terms.

RESULTS ANALYSIS

The similarity between training and testing R^2 values indicates strong generalization and no evidence of severe overfitting. Smoking status was found to have the highest positive coefficient, confirming its major impact on insurance costs. Age and BMI also positively influenced insurance charges, while sex and region had relatively smaller effects. The model successfully captures the primary real-world relationships affecting medical insurance costs. Although further improvements could be achieved using advanced models or transformations, the linear regression model provides reliable, interpretable, and meaningful predictions.

VARIABLES USED

The final processed dataset used in this study contains multiple health and demographic variables. The regression analysis primarily focuses on predicting medical insurance charges using measurable personal and lifestyle-related factors.

The dependent variable:

1. **Charges:** Charges is a continuous numerical variable representing the annual medical insurance cost of an individual. It is the target (output) variable that the regression models attempt to predict using explanatory features.

The independent variables are:

1. **Age**: Age is a continuous numerical variable representing the age of the individual in years. It is expected to have a positive relationship with insurance charges, as medical risks generally increase with age.
2. **BMI (Body Mass Index)**: BMI is a continuous numerical variable measuring body fat based on height and weight. Higher BMI values may indicate increased health risks and can lead to higher insurance charges.
3. **Smoker**: Smoker is a binary categorical variable encoded numerically:
4. **Sex:** Sex is a binary categorical variable encoded as:
5. **Region**: Region is a categorical variable representing the residential area of the individual (e.g., southeast, southwest, northeast, northwest). It is encoded numerically or using dummy variables to enable model training.

Together, these variables enable the model to analyse how demographic and lifestyle factors influence medical insurance charges and help identify the most impactful contributors, such as smoking and BMI.

DUMMY VARIABLE EXPLANATION

In regression analysis, machine learning models cannot directly interpret categorical (text-based) variables such as *sex*, *smoker status*, or *region*. Therefore, these variables must be converted into numerical format before training the model. This conversion is done using **binary encoding (dummy variables)** or **label encoding**.

A dummy variable is a numerical variable that takes only two possible values:

- **1** = Indicates the presence of a specific category
- **0** = Indicates the absence of that category

In this project, the following categorical variables were converted:

1. Smoker

The variable *smoker* originally had two categories: “yes” and “no”.

It was encoded as:

This allows the regression model to measure how smoking status affects insurance charges while keeping other variables constant.

2. Sex

The variable *sex* was encoded as:

- 1 = Male
- 0 = Female

This helps the model evaluate whether gender has any statistical influence on medical insurance charges.

3. Region

The *region* variable contains multiple categories (e.g., northeast, northwest, southeast, southwest). Since regression models cannot work with text categories directly, the variable was converted into numerical form (either through label encoding or dummy variables).

If dummy variables are used, one category is treated as the **reference category**, and the remaining categories are compared against it. This prevents multi-collinearity and ensures accurate coefficient interpretation. Using dummy variables ensures:

- The machine learning model can process categorical information
- The relationship between categories and insurance charges can be quantified
- The regression model produces meaningful and interpretable results

Thus, dummy variable encoding plays a crucial role in building an accurate Healthcare Cost Prediction model.

GITHUB IMPLEMENTATION

All code and documentation for this project are managed via GitHub. (Repository: [placeholder for repo URL]). Key milestones such as data pre-processing, model training, and evaluation have been committed with descriptive messages (e.g. *"Added Linear Regression training code"*, *"Preprocessed data and handled missing values"*). We also use version control branches for major changes. Below is a placeholder screenshot of the GitHub repository page illustrating the commit history and main code files (the actual repository link and screenshot would be provided).

Placeholder: Screenshot of GitHub repository commits and structure would appear here.

(Include screenshots of commits, GitHub UI, etc. in the final report.)

MODEL EXPLANATION

The primary model used in this project is **Linear Regression**, a supervised machine learning algorithm used to predict continuous numerical values. Since the target variable in this project, *insurance charges*, is a continuous value, linear regression was selected as an appropriate and interpretable modelling technique. The model assumes a linear relationship between the dependent variable (charges) and independent variables such as age, BMI, smoking status, sex, region, and number of children.

Before model training, categorical variables such as smoker, sex, and region were converted into numerical form using encoding techniques. The dataset was then split into training and testing subsets in an 80:20 ratio to evaluate the model's generalization ability. The Linear Regression model was trained using the training dataset, where it learned the optimal coefficients (weights) for each feature by minimizing the sum of squared errors between predicted and actual charges.

Mathematically, the model can be represented as:

$$\text{Charges} = \beta_0 + \beta_1 (\text{Age}) + \beta_2 (\text{BMI}) + \beta_3 (\text{Smoker}) + \beta_4 (\text{Sex}) + \beta_5 (\text{Children}) + \beta_6 (\text{Region}) + \varepsilon$$

Here, β_0 represents the intercept, β_1 to β_6 represent the coefficients of respective features, and ε represents the error term. Each coefficient indicates how much the predicted insurance cost changes with a one-unit increase in the corresponding variable, keeping other variables constant.

From the results obtained, the model achieved an R^2 score of approximately 0.748 on both training and testing datasets. This means that around 74.8% of the variability in insurance charges is explained by the model. The similarity between training and testing R^2 values indicates good generalization and no significant overfitting. The coefficient values revealed that smoking status has the highest positive impact on insurance costs, followed by age and BMI, which aligns with real-world healthcare trends.

Overall, the Linear Regression model provides reliable, interpretable, and meaningful predictions, making it a suitable choice for healthcare cost estimation in this project. The most significant factor influencing insurance charges is **smoking status**. From correlation analysis and regression results:

- Smoking has the **strongest positive correlation** with charges.
- The regression coefficient for *smoker* is the highest.
- Smokers have dramatically higher predicted insurance costs compared to non-smokers.
- This aligns with real-world medical evidence, since smoking increases the risk of chronic diseases and healthcare expenses.

IS UNEXPLAINED GAP PRESENT?

Yes,

An unexplained gap is present in the model. Although the Linear Regression model achieved an R^2 score of approximately 0.748, this means that about 74.8% of the variation in insurance charges is explained by the independent variables included in the model. The remaining 25.2% of the variation is not explained by the model. This portion represents the unexplained gap, also known as residual variance.

The unexplained gap exists because not all factors influencing medical insurance costs are included in the dataset. Important variables such as medical history, genetic conditions, detailed lifestyle habits (diet, physical activity), stress levels, and occupation are not available in the data. Additionally, linear regression assumes a linear relationship between features and the target variable, whereas real-world healthcare costs may follow more complex or non-linear patterns. Random variability and external factors also contribute to prediction differences.

Therefore, while the model performs well and explains a significant portion of the variation, some unexplained gap remains due to missing variables, model assumptions, and inherent randomness in medical cost data.

CODE SNIPPET

In this project, machine learning techniques were applied to predict medical insurance charges based on demographic and health-related factors such as age, BMI, and smoking status. To successfully build and evaluate the predictive model, Python programming language was used along with several powerful data science libraries. These libraries provide efficient tools for data manipulation, visualization, statistical analysis, and machine learning model development. The use of well-established libraries ensures accuracy, reliability, and scalability of the implemented solution.

```
#Import Libraries import pandas as pd import numpy as
np from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression from
sklearn.metrics import r2_score, mean_squared_error

#Load Dataset df =
pd.read_csv("insurance.csv")

#Check Missing Values rint
(df.isnull().sum())

#Encode Categorical Variables df['smoker'] =
df['smoker'].map({'yes': 1, 'no': 0}) df['sex'] =
df['sex'].map({'male': 1, 'female': 0}) df = pd.get_dummies(df,
columns=['region'], drop_first=True)

#Define Features and Target X
= df.drop("charges", axis=1) y
= df["charges"]

#Train-Test Split (80/20)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)

#Train Linear Regression Model
regression = LinearRegression()
regression.fit(X_train, y_train)
#Make Predictions y_pred =
regression.predict(X_test)
```

```
#Model Evaluation r2 = r2_score (y_test,
y_pred) mse = mean_squared_error(y_test,
y_pred)

print("R2 Score:", r2) print("Mean
Squared Error:", mse)


#View Coefficients coefficients
= pd.DataFrame({
   "Feature": X.columns,
   "Coefficient": regression.coef_}) print(coefficients)
```

This matches your project workflow and produces the $R^2 \approx 0.748$ result.

LIBRARIES USED AND EXPLANATION

1. Pandas:

Pandas is a powerful data manipulation and analysis library in Python. It was used to load the dataset (insurance.csv), explore its structure, check for missing values, remove duplicates, and perform data preprocessing tasks. Pandas provides DataFrame structures, which make it easy to filter, transform, and analyze tabular data efficiently.

2. NumPy:

NumPy (Numerical Python) is a library used for numerical computations and mathematical operations. It supports large, multi-dimensional arrays and matrices, along with high-level mathematical functions. In this project, NumPy assists in numerical processing, handling arrays, and supporting calculations required during model training and evaluation.

3. Matplotlib:

Matplotlib is a data visualization library used for generating plots and graphical representations. It was used to visualize distributions, generate scatter plots, and display model

predictions. Visualization helps in understanding trends, patterns, and relationships among variables.

4. Seaborn:

Seaborn is built on top of Matplotlib and provides advanced statistical visualizations. In this project, it was used to generate the correlation heatmap, which helps identify relationships between variables such as smoking status and insurance charges. The heatmap visually highlights strong and weak correlations.

5. Scikit-learn:

Scikit-learn is a machine learning library used for building and evaluating predictive models. In this project, it was used for:

- Splitting the dataset into training and testing sets (train_test_split)
- Implementing the Linear Regression model
- Calculating evaluation metrics such as R^2 score and Mean Squared Error

Scikit-learn provides simple and efficient tools for supervised learning tasks, making it ideal for regression-based problems.

These libraries collectively enabled efficient data preprocessing, visualization, model training, and evaluation. By integrating Pandas, NumPy, Matplotlib, Seaborn, and Scikit-learn, the project was able to build an accurate and interpretable Healthcare Cost Prediction system.

In this code:

- `Charges` is the dependent variable.
- `Age, BMI, Smoker, Sex, Children, Region (encoded)` are the independent variables.
- `Smoker (0 or 1), Sex (0 or 1), region_northwest, region_southeast, region_southwest` (*One region is dropped to avoid multicollinearity*) are included ad dummy variables.
- Healthcare cost is evaluated using **R^2, MSE, and coefficient interpretation.**

LIMITATIONS

Although the Healthcare Cost Predictor model performs reasonably well, there are several limitations in the study. First, the dataset contains only a limited number of features such as age, BMI, smoking status, sex, children, and region. Important factors that influence healthcare costs, such as medical history, genetic conditions, detailed lifestyle habits, type of insurance plan, income level, and occupation, are not included in the dataset. The absence of these variables contributes to the unexplained variance in the model.

Second, the model assumes a linear relationship between the independent variables and insurance charges. However, real-world medical costs may follow non-linear patterns or involve complex interactions between variables. Since interaction terms and non-linear transformations were not extensively implemented, some relationships may not be fully captured.

Third limitation is the right-skewed distribution of insurance charges, which indicates the presence of high-cost outliers. Although the dataset was retained to preserve real-world variability, these extreme values may influence regression performance. Additionally, the dataset consists of only 1,338 records, which may limit the model's ability to generalize to larger or more diverse populations.

Finally, the analysis is based on a publicly available dataset that may be simplified or partially synthetic for educational purposes. Therefore, the model may not perfectly reflect real insurance industry pricing mechanisms. Despite these limitations, the project provides valuable insights into healthcare cost prediction and demonstrates the effective application of machine learning techniques.

CHALLENGES FACED AND SOLUTIONS

During the project, several challenges arose.

- **Data Understanding:** Initially, I was unsure how to handle categorical features. I resolved this by researching encoding techniques and deciding on simple numeric encoding for this linear model.
- **Skewed Target:** The insurance charges were highly skewed (a few very high costs). This threatened to bias the regression. I addressed it by considering a log transformation for the target, which improved model stability during experimentation (though the final model used raw costs for easier interpretation).
- **Model Performance:** Getting good R^2 took some tuning. I tried including different feature subsets and found that age, BMI, and smoker explained most variance; adding more variables gave diminishing returns.
- **Software/Dependencies:** Installing and configuring the Python ML libraries (pandas, scikit-learn, etc.) took some troubleshooting on the first day. I overcame this by carefully following documentation. In all cases, the solution involved research (reading documentation or forums) and iterative testing.

ETHICAL CONSIDERATIONS

Predictive models in healthcare involve significant ethical responsibilities. One of the primary concerns is data privacy. Although the dataset used in this project is anonymized and publicly available, real-world healthcare data often contains sensitive personal and medical information. Such data must be handled with strict confidentiality, secure storage mechanisms, and proper access control. In practical deployment, patient consent must be obtained, and compliance with healthcare data protection regulations must be ensured. All personally identifiable information should be removed to protect patient identity and prevent misuse.

Another important concern is bias and fairness. Machine learning models can unintentionally learn patterns of bias present in historical data. For example, if certain demographic groups are underrepresented or historically underserved, the model may produce biased predictions that disadvantage those groups. In the context of healthcare cost prediction, an imbalanced dataset could lead to inaccurate cost estimation for specific populations. Therefore, it is essential to conduct bias testing, use diverse datasets, and continuously validate fairness to prevent disproportionate harm to vulnerable groups.

Transparency is also crucial in ethical model deployment. Users should clearly understand that the model provides estimated predictions based on available data and does not account for all possible medical conditions or external factors. Communicating the model's limitations ensures responsible usage and avoids misleading interpretations. Overall, maintaining privacy, ensuring fairness, and promoting transparency are essential when developing and deploying healthcare predictive models.

FUTURE ENHANCEMENTS

Although the current model performs effectively, several improvements can enhance its accuracy and real-world applicability. One potential enhancement is feature expansion. Incorporating additional health-related variables such as blood pressure, cholesterol levels, occupation, lifestyle habits, or medical history could significantly improve predictive performance. Interaction terms, such as age × smoker, could also be added to capture more complex relationships.

Another improvement involves model upgrades. Advanced machine learning algorithms such as Decision Trees, Random Forest, or Gradient Boosting models could be implemented to capture non-linear patterns in the data. Hyperparameter tuning techniques such as grid search and cross-validation could further optimize performance. From a practical perspective, developing a user interface would make the tool more accessible. A simple web application or dashboard could allow users to input their details and instantly receive cost predictions.

Additionally, deploying the model using containerization tools such as Docker and hosting it on a cloud platform would enable real-time use. Expanding the dataset to include larger and more diverse populations would also enhance the model's robustness and generalizability.

INDUSTRY USE CASES / RELEVANCE

Predictive healthcare cost models have significant applications across industries. In insurance underwriting, companies can use such models to personalize premium pricing and identify high-risk customers. Predictive analytics enhances pricing strategies and improves risk

assessment efficiency. The model developed in this project directly supports demographicbased premium estimation and smarter risk evaluation.

Healthcare providers can utilize cost predictions for resource allocation and financial planning. For example, hospitals may prioritize preventive programs for patients predicted to have higher future medical costs, such as smoking cessation initiatives. Predictive models can also contribute to anomaly detection by identifying cases where actual medical expenses significantly deviate from estimated values, potentially supporting fraud detection processes.

Public health organizations can use aggregate predictions to estimate the financial burden of risk factors such as obesity and smoking on healthcare systems. Overall, cost prediction enables proactive decision-making and strategic planning. Predictive analytics is increasingly becoming central to efficiency and competitive advantage in the insurance and healthcare industries, and this project demonstrates a foundational step toward that capability.

CONCLUSION

The Healthcare Cost Predictor project successfully applied machine learning techniques to estimate annual medical insurance charges based on demographic and lifestyle-related factors. Using the Medical Insurance Personal Dataset, the project followed a systematic approach that included data acquisition, pre-processing, exploratory data analysis, feature engineering, model training, and evaluation. After analysing the dataset, smoking status, age, and BMI emerged as the most influential predictors of healthcare cost.

A Linear Regression model was implemented as the primary predictive model due to its simplicity and interpretability. The model achieved an R^2 score of approximately 0.748 on both training and testing datasets, indicating that nearly 74.8% of the variance in insurance charges is explained by the selected features. The similar training and testing performance suggests good generalization and minimal overfitting. The results confirmed that smoking status has the most significant impact on insurance premiums, followed by age and BMI, aligning with realworld healthcare insights.Although the model performs well, an unexplained portion of variance remains due to factors not included in the dataset, such as medical history and genetic conditions. Nevertheless, the developed model provides reliable and interpretable predictions,

demonstrating the practical application of machine learning in healthcare analytics. Overall, the project highlights how data-driven approaches can improve understanding of medical cost determinants and support informed financial and health-related decision-making.

```
••• Model Performance Metrics
    --------------------------
    MAE  : 8924.62
    RMSE : 11375.26
    R²   : 0.1569
```

REFERENCES

1. Mustafa, M. Z., et al. (2025). *Patient perspective on predictive models in healthcare: translation into practice, ethical implications and limitations?* BMJ Health & Care Informatics. This paper discusses the benefits, biases, ethical implications, and limitations of predictive healthcare models.
2. United Nations. (n.d.). *Global Goals – Goal 3: Good Health and Well-being.* Official United Nations website outlining Sustainable Development Goal 3, including target 3.8 on universal health coverage and financial protection.
3. Dataquest. (2025). *Project Tutorial: Predicting Insurance Costs with Linear Regression.* An online tutorial demonstrating exploratory data analysis and linear regression techniques using an insurance dataset.
4. GeeksforGeeks. (2026). *Linear Regression in Machine Learning.* Provides definitions, working principles, assumptions, and limitations of linear regression models.
5. NVIDIA Developer Blog. (2026). *Regression Metrics: R^2 and MSE.* Explains regression evaluation metrics such as R^2 and Mean Squared Error and their significance in model performance evaluation.
6. Analytics Vidhya. (2025). *Train/Validation/Test Split – Avoid Overfitting.* Explains the importance of splitting datasets to ensure model generalization and prevent overfitting.
7. Research.com. (2026). *The Andragogy Approach: Knowles' Adult Learning Theory.* Explains principles of adult learning, including self-direction and internal motivation in knowledge acquisition.